

Implementing ANNs with TensorFlow

Session 10 - Word Embeddings

Agenda

1. Motivation
2. Statistical Language Model
3. Word Embeddings
4. Continuous Bag of Words
5. Skip Gram
6. Bonus: Seq2Seq, Attention, Transformers

Motivation

- Character-level language models can capture superficial features of the language they are trained on.
- But they fail to capture the **semantic** and **syntactic structure**



Not available due to copyright issues.

Statistical Language Model

Statistical Language Model

- Just as a character-level model a statistical language model is about learning the probabilities of sequences of words.
- These probabilities can be very useful for many purposes.
- E.g. knowing

$P(\text{'The cat climbed a tree.'}) \gg P(\text{'The cat climbed a brie.'})$


might be helpful for speech recognition.

Applications

- Speech recognition.
- Machine translation.
- Handwritten text recognition.
- Keyboards on smartphones predicting the next word.

Statistical Language Models

- What is the probability of a sentence?
- Sentence = Sequence of words.

Chain rule 

$$\begin{aligned} P(s) &= P(w_1, w_2, \dots, w_n) \\ &= P(w_n | w_{n-1}, \dots, w_1) \\ &\quad \cdot P(w_{n-1} | w_{n-2}, \dots, w_1) \\ &\quad \dots P(w_2 | w_1) \cdot P(w_1) \\ &= \prod_{i=1}^n P(w_i | w_{i-1}, \dots, w_1) \end{aligned}$$

Statistical Language Models

- Just as the character-level model a language model is always based on a corpus.
- The probabilities $P(w_i | w_{i-1}, \dots, w_1)$ can be computed by simple counting.
- The conditional probability of a word given a sequence of words is given by number of times it occurs after this specific sequence of words.

$$P(w_i | w_{i-1}, \dots, w_1) = \frac{\text{count}(w_1, \dots, w_{i-1}, w_i)}{\text{count}(w_1, \dots, w_{i-1})}$$

N-Gram Models

- An n-gram model is an approximation of that model.
- It assumes that a word is only dependent on the the $n - 1$ previous words.

$$\begin{aligned} P(s) &= P(w_1, w_2, \dots, w_n) \\ &= \prod_{i=1}^n P(w_i | w_{i-1}, \dots, w_1) \\ &\approx \prod_{i=1}^n P(w_i | w_{i-1}, \dots, w_{i-(n-1)}) \end{aligned}$$

Problem

- There is a problem with these statistical approaches:
- Let's say we find the following sentence in our text corpus:

Paris is the capital of France.

- Knowing that this sentence is valid would not have any influence on a sentence like:

Rome is the capital of Italy.

- These models are not good in generalizing to unseen sentences/combinations of words.

Problems

- The reason is that these models have no means of understanding whether two words are similar to each other (e.g. *Paris* and *Rome*, *France* and *Italy*)
- In mathematical terms this is analogous to a **one-hot encoding** (= all words are equally dissimilar).

- E.g.

$$Paris \hat{=} \begin{pmatrix} 0 \\ 0 \\ 1 \\ \vdots \\ 0 \end{pmatrix}$$


$$Italy \hat{=} \begin{pmatrix} 0 \\ \vdots \\ 0 \\ 1 \\ 0 \end{pmatrix}$$

$$France \hat{=} \begin{pmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

Word Embeddings

Embeddings


- The solution is to embed words into a high-dimensional space in which the similarity of words is represented by their distances.



Only 2-dimensional for
purpose of visualization!

Semantic Information

- In the optimal case such an embedding would not only reflect some form of similarity, but even capture semantic information.



Rome is the same to Italy as Paris is to France.

Formalization:


$$\phi('Rome') - \phi('Italy') + \phi('France') = \phi('Paris')$$



This vector captures the concept of capital.

Syntactic Information

- Similarly syntactic information can be stored this way.



Walked is the same to walking as climbed is to climbing.

Formalization: $\underbrace{\phi('walked') - \phi('walking') + \phi('climbing')} = \phi('climbed')$

This vector captures the concept of past tense.

Word Embeddings

Can we learn such a representation?

Yes!

Continuous Bag of Words

Distributional Hypothesis

- The distributional hypothesis says, that similar words occur in similar contexts.

- E.g.

The cat climbed the tree.
The dog climbed the tree. } 'Cat' and 'dog' are similar words.


- This is quite powerful, because similarity between words can be used to establish new similarity, e.g.:

Johnny has a cat.
Sarah has a dog. } Because we know that 'cat' and 'dog' are similar words we can now infer that 'Johnny' and 'Sarah' are similar words.

Continuous Bag of Words

- In the continuous bag of words model (CBOW) we try to predict a word from its context.
- The context are the words before and after the word.

Sentence from text corpus: *The cat climbed the tree.*




Continuous Bag of Words

- In practice we are given a text corpus with a vocabulary $V = \{w_1, \dots, w_{|V|}\}$.
- The text is then represented as a sequence $\vec{x}_1, \vec{x}_2, \vec{x}_3, \dots, \vec{x}_N$, where each \vec{x}_t is a one-hot vector representing the corresponding word w_i .
- Based on a defined a context window, e.g. $c = 2$, we generate training pairs

<u>Inputs</u>	<u>Targets</u>
$(\vec{x}_1, \vec{x}_2, \vec{x}_4, \vec{x}_5)$	\vec{x}_3
$(\vec{x}_2, \vec{x}_3, \vec{x}_5, \vec{x}_6)$	\vec{x}_4
\vdots	\vdots
$(\vec{x}_3, \vec{x}_4, \vec{x}_7, \vec{x}_8)$	\vec{x}_6

Embedding

- The embedding is the high dimensional representation of the word.




$$\vec{x}_t \in \mathbb{R}^{|V|}, \vec{v}_t \in \mathbb{R}^n$$

$$\mathcal{V} \in \mathbb{R}^{|V| \times n}$$


$$\vec{v}_t = \mathcal{V} \vec{x}_t$$

Embedding - LookUp



- Effectively we can substitute the matrix multiplication with a look up.
- Assume \vec{x} represents the word w_i ; \vec{x} is thus a one-hot vector with a 1 as its i -th component.
- The result of a matrix multiplication with \mathcal{V} is thus simply the i -th row of \mathcal{V} .

CBOW - Embedding



1. Calculate embeddings for all words in the context:


$$\vec{v}_i = \mathcal{V} \vec{x}_i ,$$

$$i = t - c, \dots, t - 1, t + 1, \dots, t + c$$

2. Calculate context embedding by averaging the embeddings:

$$\vec{\hat{v}}_t = \frac{\vec{v}_{t-c} + \dots + \vec{v}_{t-1} + \vec{v}_{t+1} + \dots + \vec{v}_{t+c}}{2c}$$

CBOW - Predictions



$$\vec{\hat{v}}_t \in \mathbb{R}^n, \vec{z}_t \in \mathbb{R}^{|V|}$$

$$\mathcal{U} \in \mathbb{R}^{n \times |V|}, \vec{b} \in \mathbb{R}^{|V|}$$

3. Based on the context embedding we calculate a score for each word:

$$\vec{z}_t = \mathcal{U} \vec{\hat{v}}_t + \vec{b}$$

4. By applying softmax we turn the scores into likelihoods of how probable the different words are the actual missing word:

$$\vec{y}_t = \text{softmax}(\vec{z}_t)$$

CBOW

Given a text (in one-hot vectors): $\vec{x}_1, \dots, \vec{x}_N$.

1. Chunk text into context training pairs given a context window size c :

$$(\vec{x}_{t-c}, \dots, \vec{x}_{t-1}, \vec{x}_{t+1}, \dots, \vec{x}_{t+c}), \vec{x}_t \quad \text{for} \quad t = c + 1, \dots, N - c$$

For all training pairs:

2. Calculate embeddings for all words in the context:

$$\begin{aligned} \vec{v}_i &= \mathcal{V} \vec{x}_i, \\ i &= t - c, \dots, t - 1, t + 1, \dots, t + c \end{aligned}$$

3. Calculate context embedding by averaging the embeddings:

$$\hat{\vec{v}}_t = \frac{\vec{v}_{t-c} + \dots + \vec{v}_{t-1} + \vec{v}_{t+1} + \dots + \vec{v}_{t+c}}{2c}$$

4. Based on the context embedding we calculate a score for each word: $\vec{z}_t = \mathcal{U} \hat{\vec{v}}_t + \vec{b}$

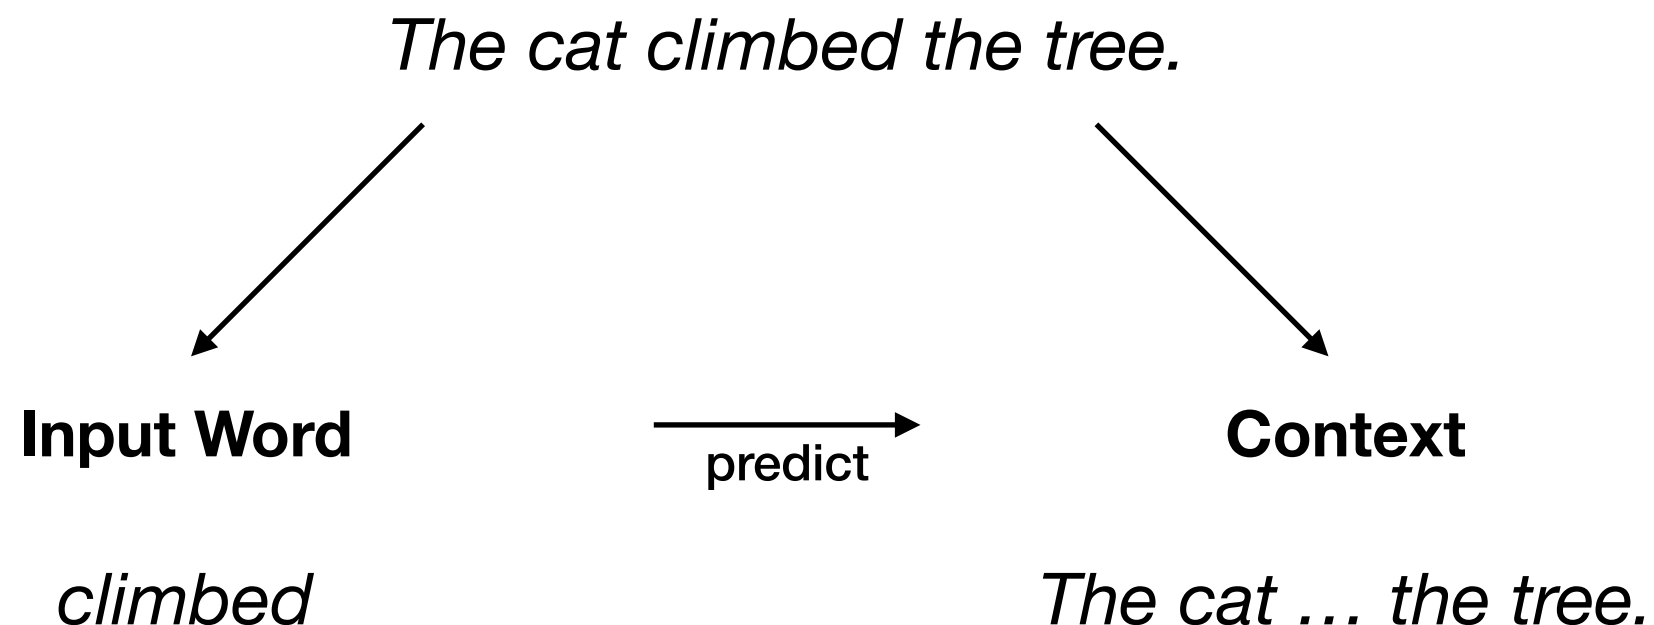
5. Apply softmax to turn score into probabilities: $\vec{y}_t = \text{softmax}(\vec{z}_t)$

6. Compute cross-entropy loss and minimize it using gradient descent.

Skip Gram

Skip Gram

- The skip gram model works as CBOW turned around.
- Now given a word we try to predict the words that occur in its context.
- This approach gives rise to more meaningful embeddings.




Skip Gram

- In practice this means that we try to predict each word in the context from the word in the center.
- A context window of size c therefore gives us $2c$ training points for each word.

Input Word	Input	Target
<i>climbed</i>	<i>climbed</i>	<i>the</i>
	<i>climbed</i>	<i>cat</i>
Context	<i>climbed</i>	<i>the</i>
<i>The cat ... the tree.</i>	<i>climbed</i>	<i>tree</i>

Skip Gram

- It uses the same basic architecture with the embedding and scoring matrix.



Skip Gram

Given a text (in one-hot vectors): $\vec{x}_1, \dots, \vec{x}_N$.

1. Chunk text into context training pairs given a context window size c :

$$\vec{x}_t, \vec{x}_{t+i} \quad \text{for} \quad t = c + 1, \dots, N - c \quad \text{and} \quad i = -c, \dots, c ; (i \neq 0)$$

For all training pairs:

2. Calculate embedding for input word: $\vec{v}_t = \mathcal{V} \vec{x}_t$
3. Based on the embedding we calculate a score for each word: $\vec{z}_t = \mathcal{U} \vec{x}_t + \vec{b}$
4. Apply softmax to turn score into probabilities: $\vec{y}_t = \text{softmax}(\vec{z}_t)$
5. Compute cross-entropy loss and minimize it using gradient descent.

Word2Vec

Word2Vec

- Skip gram was introduced in this [paper](#) by Mikolov et al. (2013).
- The model was trained on an internal google dataset (1B words, 692K unique words that occur more often than 5 times).
- The resulting embeddings can be downloaded: [link](#).
- The following are other tricks used to train the model.


Subsampling

- | | Input | Target |
|--|----------------|-------------------|
| • Some words appear very often, e.g. 'the'. | <i>climbed</i> | <u><i>the</i></u> |
| | <i>climbed</i> | <i>cat</i> |
| • These words might contain little/no information about the words around them. | <i>climbed</i> | <u><i>the</i></u> |
| | <i>climbed</i> | <i>tree</i> |
- The more often a word occurs the more often it is shown to the model.
 - To counteract this tendency we can discard a training sample based on how frequent it occurs:

Probability to discard: $P(w_i) = \max(0, 1 - \sqrt{\frac{t}{f(w_i)}})$, (e.g. $t = 10^{-5}$)

Negative Sampling

- Both CBOW and skip gram have a huge computational problem.
- They have to compute the softmax over the score vector (has length of size of vocabulary).

$$\vec{y}_t = \text{softmax}(\vec{z}_t) \quad (\vec{y}_t)^{(i)} = \frac{\exp((\vec{z}_t)^{(i)})}{\sum_j \exp((\vec{z}_t)^{(j)})}$$


This sum is the main computational load of the training and therefore really slows it down.



- Negative sampling (a variant of NCE) solves this problem.

Negative Sampling


- To understand negative sampling we have to understand what applying softmax and cross entropy loss does.
- The softmax returns the probabilities of all words to be the target word.
- Minimizing the cross entropy loss
 - maximizes the probability for the actual target word,
 - minimizes the probability for all other words.
- **Idea: Don't minimize probabilities of all other words but only of a few.**

Negative Sampling

- The problem is that by using softmax the probabilities are dependent on each other (they sum up to 1).
- So we have to find an alternative to softmax.
- Instead of applying the softmax we can apply the logistic activation function.



The values of \vec{y}_t are between 0 and 1 and sum up to 1.




The values of \vec{y}_t are just between 0 and 1.


Negative Sampling

- Now we can simply sample a few (e.g. $n = 15$) negative samples (= wrong words) and minimize their probability.

Previously:



Now:



- How to sample the negative words? Paper empirically found that $\mathcal{U}^{3/4}$ (unigram distribution (frequency) to the power of $3/4$).

Word2Vec - Example Embeddings

Not available due to copyright issues.

Bonus

Bonus

- Seq2Seq: [Blog](#), [Paper](#)
- Attention: [Blog](#)
- Transformers: [Blog](#), [Paper](#)

Any questions?

See you next week!