

# Implementing ANNs with TensorFlow

Session 13 - Practical Aspects

# Agenda

## 1. Technical

- Models
- Sequential Definition
- How to store and load a trained model?

## 2. What's going on in Research?

- Conferences
- Important Institutions and Researchers
- How to stay up to date?
- Future Research
- Courses

Model

# TensorFlow Models

- In this course we stucked with defining a model by using the class `tf.keras.layers.Layer`
- In practice it often makes things easier to initialize a model as a `tf.keras.Model`.

# tf.keras.Model

```
class Model(tf.keras.Model): ← Model class
```

```
def __init__(self):  
    super(Model, self).__init__()  
    self.conv_layer_1 = tf.keras.layers.Conv2D(  
        filters=32,  
        kernel_size=3,  
        activation=tf.keras.activations.relu,  
        input_shape=(32,32,3)  
    )  
    self.max_pool_1 = tf.keras.layers.MaxPool2D()  
    self.drop_out_1 = tf.keras.layers.Dropout(0.2)  
    self.conv_layer_2 = tf.keras.layers.Conv2D(  
        filters=64,  
        kernel_size=3,  
        activation=tf.keras.activations.relu  
    )  
    self.max_pool_2 = tf.keras.layers.MaxPool2D()  
    self.flatten = tf.keras.layers.Flatten()  
    self.fully_connected_1 = tf.keras.layers.Dense(  
        units=10,  
        activation=tf.keras.activations.softmax  
    )
```

**Everything else as usual!**

```
def call(self, x, training): ← is_training flag should be just training  
    x = self.conv_layer_1(x)  
    x = self.max_pool_1(x)  
    x = self.drop_out_1(x, training=training)  
    x = self.conv_layer_2(x)  
    x = self.max_pool_2(x)  
    x = self.flatten(x)  
    x = self.fully_connected_1(x)  
    return x
```

# Compile the Model

```
model = Model()  
optimizer = tf.keras.optimizers.Adam()  
loss = tf.keras.losses.CategoricalCrossentropy()  
accuracy = tf.keras.metrics.CategoricalAccuracy()  
model.compile(optimizer, loss, metrics=[accuracy])
```

- **Initialize the model, optimizer, loss and metrics (e.g. accuracy).**
- **You have to make sure that the loss and the accuracy can be computed on your data (e.g. you have to create the one-hot labels before).**

# Train the Model

```
(train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.cifar10.load_data()
train_labels = train_labels[:,0]
train_labels = tf.one_hot(train_labels, depth=10).numpy()
test_labels = test_labels[:,0]
test_labels = tf.one_hot(test_labels, depth=10).numpy()
train_images = np.array(train_images, dtype=np.float32)
test_images = np.array(test_images, dtype=np.float32)
train_images, test_images = train_images / 255.0, test_images / 255.0
```

**You can simply feed the data as numpy arrays!**

```
model.fit(train_images, train_labels, batch_size=32, epochs=3, validation_data=(test_images, test_labels))
```

**This single line will train the model on the corresponding training data for 3 epochs with a batch size of 32 and validate once a epoch on the validation data.**

Train on 50000 samples, validate on 10000 samples

Epoch 1/3

50000/50000 [=====] - 44s 884us/sample - loss: 1.5393 - categorical\_accuracy: 0.4370

- val\_loss: 1.3364 - val\_categorical\_accuracy: 0.5153

Epoch 2/3

50000/50000 [=====] - 45s 908us/sample - loss: 1.1922 - categorical\_accuracy: 0.5755

- val\_loss: 1.1555 - val\_categorical\_accuracy: 0.5940

Epoch 3/3

50000/50000 [=====] - 43s 863us/sample - loss: 1.0579 - categorical\_accuracy: 0.6252

- val\_loss: 1.0234 - val\_categorical\_accuracy: 0.6406

# Model Summary

You can even read out a summary of the model.

```
model.summary()
```

Model: "model\_1"

| Layer (type)                   | Output Shape | Param # |
|--------------------------------|--------------|---------|
| conv2d_3 (Conv2D)              | multiple     | 896     |
| max_pooling2d_2 (MaxPooling2D) | multiple     | 0       |
| dropout_1 (Dropout)            | multiple     | 0       |
| conv2d_4 (Conv2D)              | multiple     | 18496   |
| max_pooling2d_3 (MaxPooling2D) | multiple     | 0       |
| conv2d_5 (Conv2D)              | multiple     | 36928   |
| flatten_1 (Flatten)            | multiple     | 0       |
| dense_2 (Dense)                | multiple     | 65600   |
| dense_3 (Dense)                | multiple     | 650     |
| Total params: 122,570          |              |         |
| Trainable params: 122,570      |              |         |
| Non-trainable params: 0        |              |         |



# Sequential Model Definition

- The sequential model definition makes things even easier.

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Conv2D(filters=32,
                                   kernel_size=3,
                                   activation=tf.keras.activations.relu,
                                   input_shape=(32,32,3)))
model.add(tf.keras.layers.MaxPool2D())
model.add(tf.keras.layers.Dropout(0.2))
model.add(tf.keras.layers.Conv2D(filters=32,
                                   kernel_size=3,
                                   activation=tf.keras.activations.relu))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(units=10,
                                   activation=tf.keras.activations.softmax))
```

- These lines are all you need for the definition of the model.

# Saving a Model

- Saving a model is very easy:

```
model.save('stored_model.h5')
```

- This will store the weights plus the optimizer and its state.
- If you reload the model and train on you will start from the exact point that you stopped.

```
new_model = tf.keras.models.load_model('stored_model.h5')
```

What's happening  
in Research?

# Conferences

- In opposite to most fields of research machine learning papers are mainly published on conferences.
- It can be quite interesting to visit conferences to get into touch with researchers and get a feeling for the field.
- The most important conferences are:
  - International Conference on Learning Representations (ICLR)
  - International Conference on Machine Learning (ICML)
  - Neural Information Processing Systems (NeurIPS)
  - AAAI Conference on Artificial Intelligence
  - Application oriented ones: e.g CVPR, ICCV, ACL

# Conferences

- Sadly the most of these conference do not take place in Europe, but most of the time in the US or Canada.
- But ICML 2020 will be in Vienna!
- And in Tuebingen there will be the Machine Learning Summer School with some popular international speakers (<http://mlss.tuebingen.mpg.de/2020/speakers.html>).

# NeurIPS Challenges

- The NeurIPS conference always features challenges on specific topics (e.g. Causality for Climate Change).
- Everyone can take part in these competitions and you usually only need to submit code/results and a small report. (not a complete publication).
- These can be a nice way to take part in a conference.
- [Link](#)

# Open Review

- Sometimes conferences organize the reviewing process via <https://openreview.net>.
- ICLR does this every year.
- It can be interesting to read the reviews and the author's answers to get a feeling for how research is conducted.

# Conference Statistics

- For most conferences you will find published statistics.
- These give you information about the institutions and authors and thus show you, where most research is happening.
- In the last years the contribution of industrial institutions increased rapidly.



# Industrial Research

- Google (incl. Google Brain, Google Deep Mind)
- All other big tech companies: Microsoft, Facebook, IBM, Amazon
- Tons of other bigger companies e.g.: Bosch, Adobe, NVIDIA
- Smaller startup AI companies e.g.: CuriousAI, CriteoAI, Vicarious
- OpenAI: quite few but high quality publications

# Academia

- Main academic players in U.S. and Canada:
  - MIT (Tomaso Poggio, Josh Tenenbaum)
  - Berkeley (California) (Sergey Levine, Pieter Abbeel)
  - Stanford (Li Fei-Fei, Chris Manning)
  - Carnegie Mellon (Russ Salakhutdinov)
  - Georgia Tech
  - Toronto (Vector Institute + Hinton)
  - Montreal (MILA + Bengio)

# Academia - Europe

- GB: U. o. Oxford, U. o. Cambridge, U. o. Edinburgh, University College London, Imperial College London
- Switzerland: ETH Zuerich, EFPL Lausanne, IDSIA Lab (Schmidhuber)
- Germany: U. o. Tübingen (Schölkopf, Bethge), KIT (Karlsruhe), TU Darmstadt, U. o. Saarland (Saarbrücken)
- Others: Aalto University (Finland), KTH (Stockholm, Schweden), U. o. Amsterdam (Max Welling)

# Get into it

- If you want to get into the field of research and see what is going on:
  - Follow the conferences and read important papers (spotlight).
  - Follow ArXiv. But be aware that anyone can upload papers here so it's very crowded and unreviewed.
  - Twitter (twitter-active researchers who might be a good starting point: @hardmaru, @rasbt, @jeremyphoward, @fchollet, @zacharylipton)
  - Online blogs/news: Machine learning mastery, Lil-log, Deep Mind, Deep Hunt, Hacker Noon

# Courses

- Online courses: fast.ai, deeplearning.ai, Coursera
- Stanford NLP course (Chris Manning)
- Deep Mind Reinforcement Learning (David Silver)
- MIT Introduction to Deep Learning

# Hot Topics

# Artificial vs. Human Intelligence

- There is a very general discussion going on, whether deep learning alone is suitable to model actual human intelligence.
- Why? Because processing raw inputs is only achievable with deep learning, while many higher cognitive functions (e.g. counting, compositional understanding) seem easier to implement on a symbolic level.
- Advocates for Deep Learning: everyone who is doing deep learning including the “godfathers” : Y. Bengio, Y. LeCun, G. Hinton.
- Critiques from Cognitive Scientist: e.g. Gary Marcus (author of Rebooting AI), Brenden Lake

# Resources

- Building Machines that Think and Learn like People (Lake et. al, [Paper](#)) (in general publications from [B. Lake](#))
- DeepMind Comment on Building Machines that Think and Learn like People ([link](#))
- Talk of Y. Bengio on lifting Deep Learning to the next level ([link](#))
- [Review Paper](#) on reconciling deep learning and symbolic AI.



# Human-Like Category Learning

- The way how humans learn categories differs widely from how artificial neural networks are trained on a dataset of different classes.

## Neural Networks ...

- Have access to all categories from the start of training.
- Don't need to deal with images, which do not belong to any known categories.
- Need thousands of samples to learn a category in a robust way.
- Can't explain the underlying concept of a category.




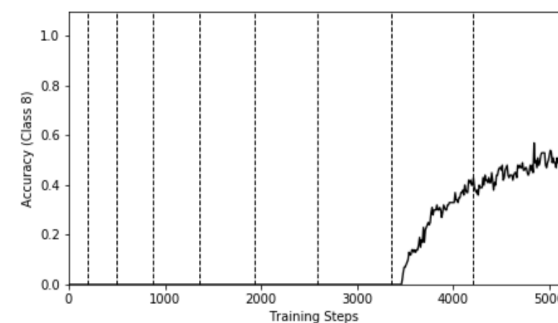
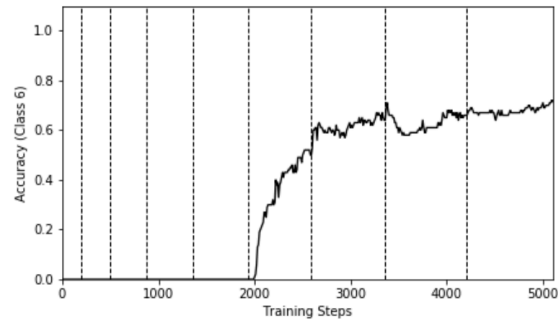


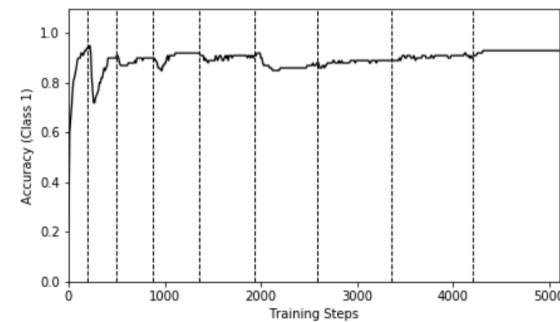
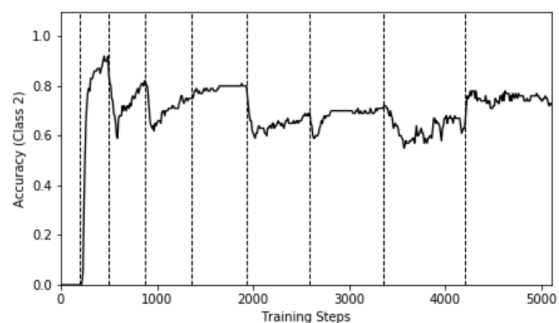

## Humans ...

- Continually learn novel categories.
- Can detect when a sample is of an unknown category.
- Can learn categories from as few as one sample.
- Are somewhat aware of the underlying concept, which defines the category.

# Continual Learning


- The task of learning categories one after the other is called **continual learning**.
- If neural networks learn this way one observes a phenomenon called **catastrophic forgetting**, which means that a network gets worst in recognizing “old” categories after it learned new categories.

## Simple Example on MNIST



# Out-of-Distribution Detection


- The task of understanding when a sample belongs to a category which is unknown is called **out-of-distribution detection** (OODD).
- Neural networks usually do not have this ability and sometimes assign high probabilities to out-of-distribution samples:



(a) Prediction: 6  
Confidence: 0.92




(b) Prediction: 9  
Confidence: 0.92



(c) Prediction: 4  
Confidence: 0.96



(d) Prediction: 9  
Confidence: 0.45



(e) Prediction: 6  
Confidence: 0.72



(f) Prediction: 9  
Confidence: 0.77

# One-/Few-Shot Learning

- The task of learning a new category from a few samples is called **one-** or **few-shot learning**.
- Humans are quite good at it while a classical neural networks performance drops heavily.
- A popular task to test on this is a so called **k-shot-view task**.

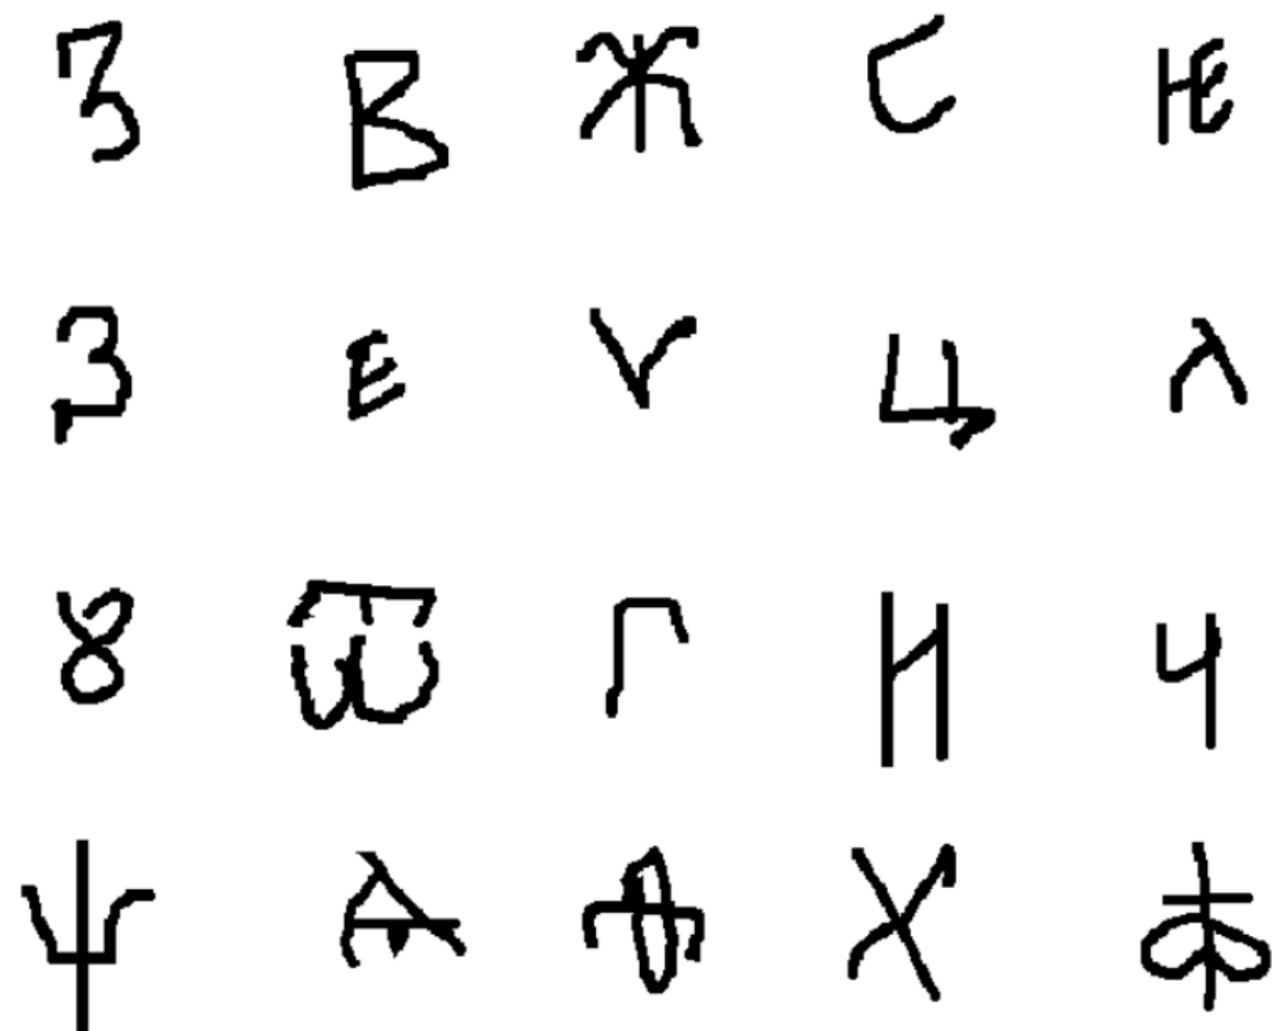
**Example: 1-shot-20-view task**

**Try to learn/remember the following character!**



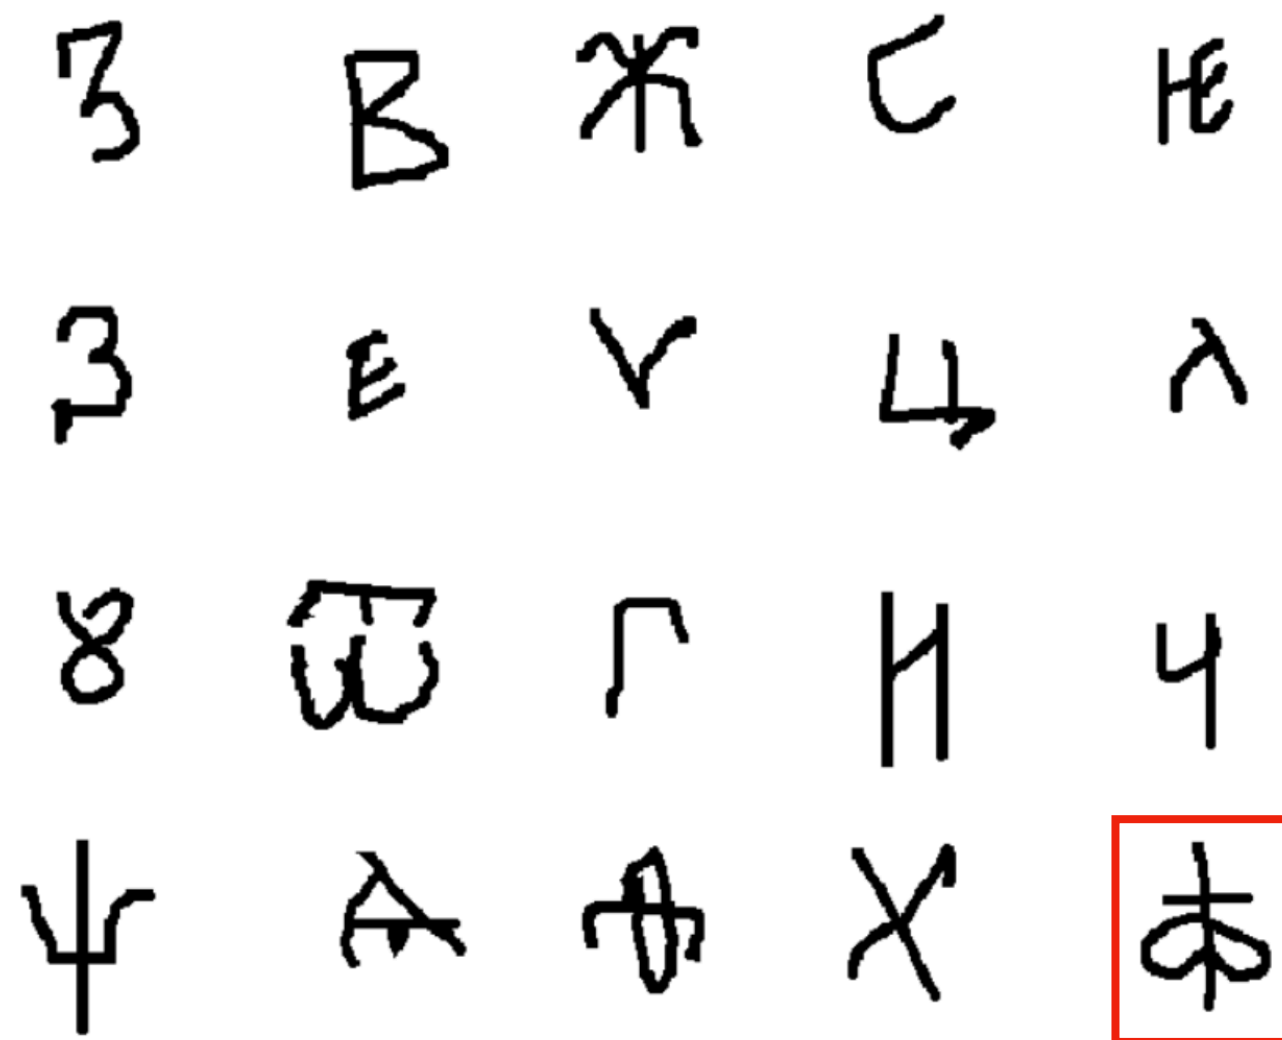
# One-/Few-Shot Learning

Which of the following is the same character?



# One-/Few-Shot Learning

Which of the following is the same character?



# Omniglot Challenge

[Paper](#)

- The 1-shot-20-view task was presented as part of the **Omniglot Challenge**, a set of tasks designed to mirror human-level concept learning abilities.

Not available due to copyright issues.

i) **One-shot learning  
(classification)**

ii) **One-shot learning  
(generation)**

iii) **Concept  
decomposition**

iv) **Generation of new  
concepts**

# Omniglot Challenge

- Since the publication of the challenge a lot of research was done on one-/few-shot learning (quite successfully for simple characters).
- Some work was done on one-shot generation, but very few was achieved on the decomposition part.
- See [review paper](#) from the authors of the omniglot challenge for a comprehensive summary of all relevant approaches.



# Compositional Understanding

- The ability to decompose a concept touches on a very general notion of **compositional understanding**.
- Compositionality is an extremely important feature of human language (the meaning of a sentence is determined by its structure and the meaning of its words).
- Neural networks fail to leverage this aspect in their learning process.

# Compositional Understanding

[Paper](#)

- Example Task for few-shot learning of compositional instructions:

Not available due to copyright issues.

# Explainable AI

- The task of decomposing a concept into its “components” also touches on the field of **explainable AI**.
- Explainable AI is an important line of research working on ways of extracting human-interpretable rules from machine/deep learning algorithms.

# Safety and Ethics in AI

- Explainable AI itself is part of a broader line of research summarized under **Safety of AI** ([Stanford Research Center](#)).
- This includes AI ethics (e.g. autonomous driving/ weapons), building trustworthy AI, building robust AI.
- This field is extremely important and will presumably get more and more important over the next years.

# Other Important Fields

- Reinforcement Learning (incl. for robotics) - lecture 12
- Unsupervised Learning of Disentangled Representations - lecture 11

# Final Project

# Goal

- The goal of the final project is to test whether you can:
  - read a research paper,
  - abstract and comprehend the most important aspects,
  - implement the architecture and experiment in TensorFlow
  - and give a clear explanation of what you did.
- Of course it is nice to have an exciting project with awesome results, but this should not be your focus.

# Submission Format

- The submission has to contain **code** and **written text**, explaining your project.
- The written text should contain:
  - Introduction (+relevant background knowledge, incl. citations).
  - Main part. Explaining the publication (architecture, loss, experiments...) accompanied with the important parts of the code.
  - Results + Visualization + Discussion
- **Choice 1**: One Jupyter notebook. With the written text in markdown and the code in the cells in between.
- **Choice 2**: Submit code via a GitHub repo. And write a pdf-report (using LaTeX). In this case include screenshots of important parts of the code.



# Assessment

- The assessment of the project will be based on:
  - The displayed understanding in the report
  - Readability of report/code (include comments)
  - Clear visualization of experimental results.
  - Effectiveness of implementation.

**It is fine to look for implementations in the web and get help and inspiration but be aware that we will scan for plagiarism, which will result in a FAIL!**

# Grid

- You can use the grid system for your final project.
- This means you can use computers located in the institute, which have GPUs.
- Working on the grid can be annoying though.
- If you would like to work on the grid please send me a mail until the end of this week.
- There is a Stud.IP group “Grid-Computing at the ...”, where you can find infos on how to use it.

# Last Words

# Feedback

- The team and I would love to get some feedback on the course.
- Please fill out the following [Google Form](#).

Goodbye!